

Sign Language

Gestures to interpret

Ketan Patel, Noah Guralnik,
Matthew Benvenuto



The ASL Sign Language Gestures

- Here is a list of all of our gestures that we are using to build an Eigen-sign Matrix.
- Each column of the Eigen-Sign Matrix will be a unique image (sign).



One Image

Analyzing the mathematical structure of an individual image

The structure after removing color and creating a vector for an image

```
The image shape is: (200, 200, 3)
The image height is: 200
The image width is: 200
The image depth(color channels) is: 3
The total pixels per image for a colored image is: 120000
The total pixels per image for a black and white image is: 40000.0 which is 3 times smaller
```

```
Pre-reshapping (200, 200, 3)
Reshaping: (120000, 1)
Reshaped, removing color: (40000, 1)
```

Mega Matrix X

The Shape of X1: (40000, 27)

- This matrix is the key to the whole design.
 - Applied for sign recognition
 - Applied for sign detection

```
#This matrix is built specifically to make the eigensign matrix where each column is a sign image in vector representation
def mega_matrix_X(all_images):
    color_weight = [0.2125, 0.7154, 0.0721]; # LUMA-REC.709
    number_samples = len(all_images)
    # total pix is defined to be new_height * new_width a few lines above ^^
    X = np.zeros((total_pixs, number_samples), dtype=np.float64) #Setting up a matrix X of column with all number_samples
    i = 0
    for image in all_images:

        #This is arithmetic to convert the image into gray scaale using color weight and build X vector
        pix_gray = np.dot(plt.imread(image)[..., 0:3], color_weight)
        Xi = pix_gray.reshape(-1,1)

        #Filling the Mega X Matrix
        if Xi.shape[0]==total_pixs:
            X[:,i]=Xi[:,0]
            i=i+1
        else:
            print("Discarding image because of incorrect dimension: {}".format(image))
    return X
```

You might be thinking what is that?

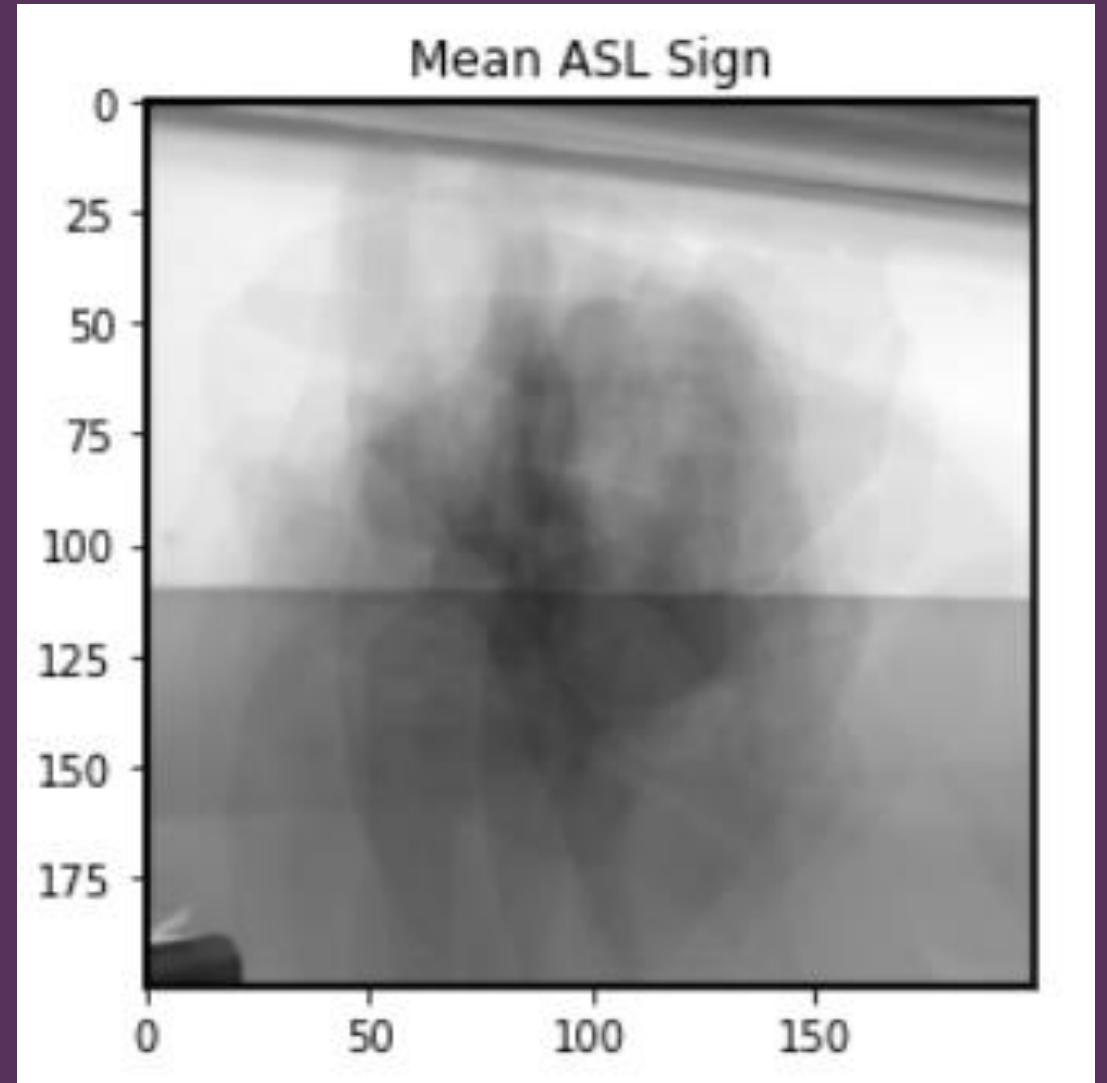
- Is that a pinky?
- Is that the Flash waving his hand?
- Is it Casper the ghost?
- Actually, this is just the Mean Sign of our ASL sign language gestures



[This Photo](#) by Unknown Author is
licensed under [CC BY-SA](#)




[This Photo](#) by Unknown Author is
licensed under [CC BY-SA](#)





Sign Detection Implemented in Python

$L2 \text{ Norm} \| (x_c - \text{mean face}) - \text{proj}(x_c \text{ on sign space}) \|$
where x_c = candidate image



- Projection of an image x on the sign space $U_{(k)} U_{(k)}^T (x - \bar{X}) = U_{(k)} U_{(k)}^T \tilde{x}$.
- Distance of the projection from the original

$$\delta = \|\tilde{x} - U_{(k)} U_{(k)}^T \tilde{x}\|$$

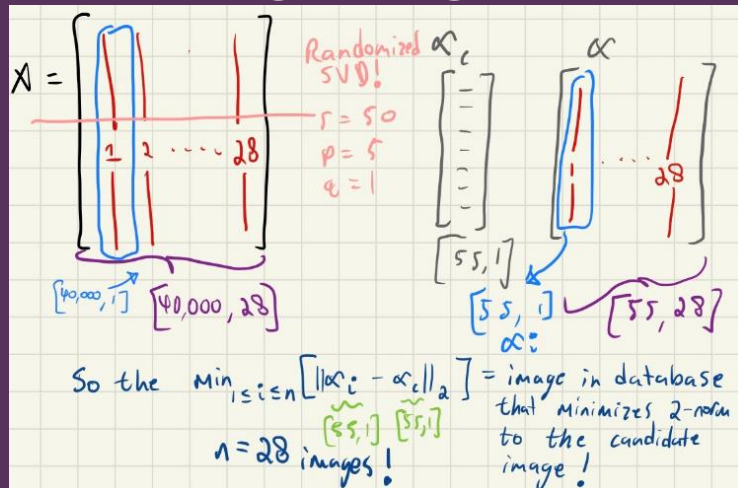
```
def sign_detection(image_file_path):
    image = Image.open(image_file_path)
    new_200 = image.resize((200, 200))
    candidate_image = np.array(new_200)
    candidate_image = np.dot(candidate_image[...], 0:3], color_weight)

    #Running the Randomized SVD: Hyperparameter choices are optimized and tuned
    r = 400 # Target rank
    q = 1  # Power iterations
    p = 5  # Oversampling parameter
    rU, rS, rVT = randomized_SVD(X1,r,q,p)
    rU_rU_T = np.diag(rU).dot(rU.T)
    #Then I will build x_tilda , and the original weight_vector
    x_tilda = candidate_image-Mean_sign
    x_tilda = x_tilda.reshape(-1,1)

    delta = norm(x_tilda-rU_rU_T.dot(x_tilda))
    return delta
```

Image Recognition

- We are applying randomized SVD to the X1 mega matrix so that we can project some candidate weight vectors and compare to the database.
- Simultaneously we are resizing, reshaping, a colored candidate image into a (200,200) image so that math is aligned.
- We are iterating through the columns of



- Find the linear coding (or weight vector) in projecting OF an image $\tilde{x} = x - \bar{x}$ to the sign space:

$$\alpha = U_{(k)}^T \tilde{x}$$

- An image could belong to a class c if the

$$\min_i (\|\alpha_i - \alpha_c\|) \text{ s.t. } i \text{ is the Eigen-Sign in the Database}$$

```
#This will take in a list of file paths, and will return a recognized sign letter
def candidate_recognition_predictor(image_file_path):
    #Reads the image file path in and then converts into gray scale
    image = Image.open(image_file_path)
    new_200 = image.resize((200, 200))
    candidate_image = np.array(new_200)
    candidate_image = np.dot(candidate_image[... , 0:3], color_weight)

    #Running the Randomized SVD: Hyperparameter choices are optimized and tuned
    r = 50 # Target rank
    q = 1 # Power iterations
    p = 5 # Oversampling parameter
    rU, rS, rVT = randomized_SVD(X1,r,q,p)
    weight_vector= np.diag(rS).dot(rVT)

    #Then I will build x_tilda , and the original weight_vector
    #x_tilda = candidate_image-Mean_sign
    #x_tilda = x_tilda.reshape(-1,1)

    x_tilda= candidate_image.reshape(-1,1)-M.reshape(-1,1)

    # Now we want to apply the projection of the alpha : Projection of the candidate weight
    alpha = rU.T.dot(x_tilda)

    #original_weight_vector : Projection of the database weights
    original_weight_vector = rU.T.dot(X1-M.reshape(-1,1))

    #Build are vector of candidate to the database
    vector = []
    length_columns = len(original_weight_vector[0])
    #Essentially subtracting the candidate matrix from the each individual column of alpha matrix
    # and appending the norm to the list
    for i in range(length_columns):
        vector.append(norm((original_weight_vector.T[i] - alpha),2))
        # vector.append(norm((original_weight_vector.T[i] - alpha),2)/norm(original_weight_vector.T[i],2))
    #print(vector)
    #Before returning we could even quickly plot the image that is closest via the 2-norm
    recognized_image_index = np.argmin(np.array(vector))
    recognized_image = imread(sign_test_data[recognized_image_index])
    plt.imshow(recognized_image,cmap=plt.cm.gray)

    #returning the index of the min vector
    return np.argmin(np.array(vector))
```


Given the Candidate Image of Matthew Benvenuto's Peace Sign



- The algorithm for 30 different runs measuring the L2 norm against the candidate image produced similar looking images but was not fully accurate!

Matt peace sign recognition

```
[191] matt_delta = sign_detection("/content/drive/MyDrive/matt_peace.JPG")  
print("Matt peace sign with white background DELTA:",matt_delta)
```

```
matt_peace = np.dot(plt.imread("/content/drive/MyDrive/matt_peace.JPG")[..., 0:3], color_weight)  
plt.imshow(matt_peace,cmap=plt.cm.gray)
```

Matt peace sign with white background DELTA: 6801.981628117551
<matplotlib.image.AxesImage at 0x7f961533f1d0>





Given the Candidate Image of Ketan Patel's Peace Sign

- The algorithm for 30 different runs measuring the L2 norm against the candidate image produced similar looking images and was super accurate recognizing the peace sign!

Ketan Peace Recognition

```
[192] ketan_hand = sign_detection("/content/drive/MyDrive/ketan_hand.JPG")  
      print("Ketan fist DELTA:",ketan_hand)  
  
      ketan_hand_img = np.dot(plt.imread("/content/drive/MyDrive/ketan_hand.JPG")[...,:3], color_weight)  
      plt.imshow(ketan_hand_img,cmap=plt.cm.gray)
```

Ketan fist DELTA: 13991.691428534876
<matplotlib.image.AxesImage at 0x7f9d1597e510>




The Fist Recognition was the best yet, not exact but accurate!

Ketan Fist Recognition

```
[37] ketan_fist = sign_detection("/content/drive/MyDrive/ketan_fist.JPG")  
     print("Ketan fist DELTA:",ketan_fist)  
  
     ketan_fist_img = np.dot(plt.imread("/content/drive/MyDrive/ketan_fist.JPG")[..., 0:3], color_weight)  
     plt.imshow(ketan_fist_img,cmap=plt.cm.gray)
```

<matplotlib.image.AxesImage at 0x7f961a3fb150>



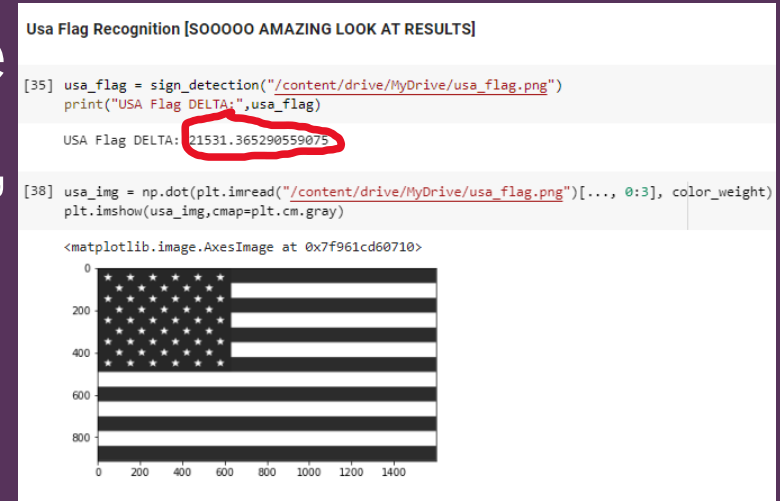
If we were to use a different candidate image like a black and white USA flag, what would you expect?



Photo by [Luke Michael](#) on [Unsplash](#)



Photo by [Nathan Fertig](#) on [Unsplash](#)



Frog Hand

- Image Recognition results make sense choosing the gesture with most curves, and the one that best represents forearm in picture.




```
import scipy.misc
from scipy import ndimage
import matplotlib.pyplot as plt
frog_hand = np.dot(plt.imread("/content/drive/MyDrive/frog_hand.JPG")[..., 0:3], color_weight)
rotated_frog_img = ndimage.rotate(frog_hand, 1.5*60)

frog_hand_delta = sign_detection("/content/drive/MyDrive/frog_hand.JPG")
print("Frog hand DELTA:", frog_hand_delta)

plt.imshow(rotated_img, cmap=plt.cm.gray)
```

```
Frog hand DELTA: 22446.35077263324
<matplotlib.image.AxesImage at 0x1961a004450>
```



Video to Image Converter: This is a cool application of this Eigen-Signs

Video to image converter: This is useful if we want to take a personal video of our hands and to detect the overall translation capabilities of our model

The design schema for our recording can be as follows:

- 1). Start the video with a ASL sign.
- 2). Wait 5 seconds and then throw up another sign.
- 3). Repeat this block style of pause and sign, until communication is finished.


If we want we can artificially upload a string of ordered images and then analyze the ordered communication accuracy.[i.e. a sentence in ASL]

Lets try recording a live video and cutting the image frame by frame! There is some accuracy but could be better


Video Frame by Frame Sliced



Image Recognition




Corresponding Candidate in Database



*The Video Model's gesture recognition did not produce the best quality images, but even still the gesture model managed pretty well, **bottom two rows are close , and third column nice***

2800 images



Not so good mean sign,
cause for bad results!

```
X1 = mega_matrix_X(all_images)
print(X1.shape)
```

```
(40000, 2800)
```

```
print("There are {} images".format(len(all_images)))
```

```
There are 2800 images
```

```
base_path_train = "C:/Users/Matthew Benvenuto/Downloads/archive/asl_alphabet_train/asl_alphabet_train/"
train_dir=os.listdir(base_path_train)[: ]
print(len(train_dir))
print(train_dir)
```

```
file_paths = []
sign_images_train = []
for subfolder in train_dir:
    if(subfolder == "nothing"): #not including nothing directory, removing from list because contributes to mean face making bad
        continue
    file_paths = [base_path_train + '/' + subfolder]
    sign_images_train = sign_images_train + file_paths
```

```
new_file_paths = []
all_images = []
dictionary= {} # we need to use this if we want to do any model classification!
num_images = 0

for subfolder in sign_images_train:
    subfolder_dir = os.listdir(subfolder)
    count = 0
    for i in subfolder_dir:
        # Collecting the first 200 images in each file !
        if(count<100): ##Might be able to put an if statement here to collect a smaller amount of images
            new_file_paths = [subfolder + '/' + i]
            all_images = all_images + new_file_paths
            num_images = num_images+1
        else:
            continue
    count = count+1
```


Future Directions

- Using more data points, 2800 training gestures and getting a more accurate mean sign!
- Optimizing the video to image frame converter so that the video stream can recognize gestures more precisely.
- Using adaptive thresholding, edge detection, transforming aligned images, blurring background, and hand detection to control the variables effecting the errors.

Machine Learning Model Architecture: After 5 epochs- 93% accuracy

```
data_dir = base_path_train
target_size = (64, 64)
target_dims = (64, 64, 3) # add channel for RGB
n_classes = 29
val_frac = 0.1
batch_size = 64

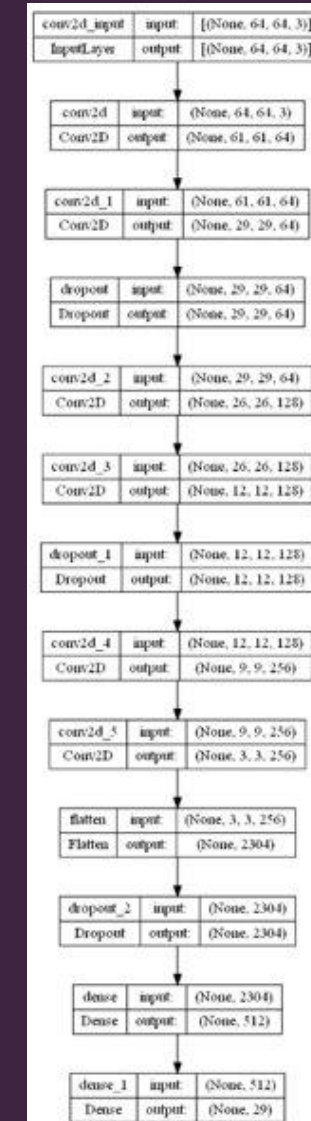
data_augmentor = ImageDataGenerator(samplewise_center=True,
                                     samplewise_std_normalization=True,
                                     validation_split=val_frac)

train_generator = data_augmentor.flow_from_directory(data_dir, target_size=target_size, batch_size=batch_size, shuffle=True, subset="training")
val_generator = data_augmentor.flow_from_directory(data_dir, target_size=target_size, batch_size=batch_size, subset="validation")
```

```
my_model.fit_generator(train_generator, epochs=5, validation_data=val_generator)

C:\Users\Matthew Benvenuto\AppData\Local\Temp\ipykernel_14940\2957552873.py:1: UserWarning: 'Model.fit_generator' is deprecated and will be removed in a future version. Please use 'Model.fit', which supports generators.
  my_model.fit_generator(train_generator, epochs=5, validation_data=val_generator)

Epoch 1/5
1224/1224 [=====] - 639s 521ms/step - loss: 2.0001 - accuracy: 0.3688 - val_loss: 1.2687 - val_accuracy: 0.5831
Epoch 2/5
1224/1224 [=====] - 488s 399ms/step - loss: 0.6367 - accuracy: 0.7825 - val_loss: 0.9283 - val_accuracy: 0.6961
Epoch 3/5
1224/1224 [=====] - 493s 402ms/step - loss: 0.3462 - accuracy: 0.8818 - val_loss: 0.7474 - val_accuracy: 0.7794
Epoch 4/5
1224/1224 [=====] - 484s 395ms/step - loss: 0.2440 - accuracy: 0.9175 - val_loss: 0.6359 - val_accuracy: 0.8226
Epoch 5/5
1224/1224 [=====] - 496s 405ms/step - loss: 0.1885 - accuracy: 0.9374 - val_loss: 0.7036 - val_accuracy: 0.8148
```



Work Cited

- All images have been properly cited within the slides and below:
 - [This Photo](#) by Unknown Author is licensed under [CC BY-SA-NC](#)
 - Photo by [Luke Michael](#) on [Unsplash](#)
 - Photo by [Nathan Fertig](#) on [Unsplash](#)